

50-DAY CHALLENGE · PHASE 2

Claude Code CLI — Complete Guide

From installing the terminal tool to orchestrating multi-agent workflows — how it connects, how it works, how to use it effectively, plus a full interview-prep round and a one-page big-picture map.

"An LLM in a chat box answers questions. Claude Code lives inside your project — it reads your files, runs your commands, and edits your code. You stop copy-pasting and start delegating."

15

PARTS

11

DIAGRAMS

43

INTERVIEW Qs

Table of Contents

- Why this phase exists
 - Part A — What is Claude Code?
 - A1. From chatbot to coding agent · A2. CLI, why the terminal? · A3. Assistant vs Agent · A4. Where Claude Code runs · A5. Key terms glossary
 - Part B — Getting Connected
 - B1. Install · B2. Authentication · B3. Your first session · B4. The interface & modes
 - Part C — How It Works Internally
 - C1. The agentic loop · C2. The context window · C3. Tools · C4. Permissions
 - Part D — Models
 - D1. The model family · D2. How they differ · D3. Switching models · D4. Cost & speed
 - Part E — CLAUDE.md (Project Memory)
 - Part F — Commands & Controls
 - Part G — Using Claude Code Effectively
 - Part H — MCP (Model Context Protocol)
 - Part I — Agents & Workflows
 - Part J — Real-World Scenarios
 - Part K — Exercises & Deliverables
 - Part L — Q&A for Every Topic
 - Part M — Interview Preparation
 - R1 Fundamentals · R2 Internals · R3 Models · R4 CLAUDE.md & Commands · R5 MCP · R6 Agents · R7 Product Differences · R8 Scenarios · R9 Behavioural · R10 Rapid-Fire
 - Part N — Resources & Quick Reference
 - Part O — The Big Picture: How It All Connects
-

Why this phase exists

Phase 1 taught you to *talk* to an AI (prompt engineering). Phase 2 teaches you to *work with* one inside a real codebase.

The difference is enormous. A chat model can only see the text you paste. Claude Code can open your files, search your project, run your tests, fix the failures, and commit the result — all from your terminal, while you supervise. This is the tool professional developers now use daily, and the skill that separates "I use AI to write snippets" from "I ship features with AI."

By the end you will be able to: install and connect Claude Code, pick the right model for a job, write a `CLAUDE.md` that makes it smarter about *your* project, connect external tools with MCP, and orchestrate multiple agents to do work in parallel.

Part A — What is Claude Code?

Start from zero. If you finished Phase 1, you already know what an LLM is — that's all the prior knowledge you need.

| A1. From chatbot to coding agent

Simple definition: Claude Code is a command-line tool (a program you run in your terminal) that puts Claude *inside* your project folder. Instead of you copying code into a chat window and pasting answers back, Claude reads and edits the files directly.

Analogy: A chatbot is a brilliant consultant on the phone — you describe your problem, they describe a solution, and *you* do all the typing. Claude Code is that same consultant sitting at your desk, hands on your keyboard, with permission to open files and run commands while you watch. The intelligence is the same; the *reach* is completely different.

What changes in practice:

- **Chatbot:** "How do I add a dark-mode toggle?" → it explains → you implement it yourself across 4 files.
- **Claude Code:** "Add a dark-mode toggle" → it *finds* the 4 files, edits them, runs the app to check, and shows you the diff.

Mindset takeaway: Claude Code is not a smarter chatbot. It is an *agent* — a model that can take actions in the real world (your filesystem, your terminal) in a loop until the job is done.

| A2. CLI, why the terminal?

CLI = Command-Line Interface — a text-based way to control a program by typing commands, instead of clicking buttons (that's a GUI, Graphical User Interface).

Why did Anthropic build a *terminal* tool instead of a fancy app? Because the terminal is where code already lives:

- Your project files are in folders the terminal can reach.
- Your build tools, tests, and Git already run as terminal commands.
- Giving Claude a terminal means giving it the *same tools you use* — no special integration needed.

Analogy: A GUI app is like a restaurant with a fixed menu — you can only order what the buttons allow. The terminal is like a kitchen — anything you can cook, Claude can cook too, because it has the same utensils.


Don't panic about the terminal. You type in plain English, not cryptic commands. "Fix the login bug" is a valid instruction. The terminal is just the *doorway*; the conversation inside is normal language.

| A3. Assistant vs Agent (the key shift)

This is the single most important idea in Phase 2.

	Assistant (chatbot)	Agent (Claude Code)
Sees	Only what you paste	Your whole project, on demand
Does	Produces text	Takes actions: read, edit, run, search
Loop	One question → one answer	Keeps going until the goal is met
You	Do the work it describes	Supervise the work it performs
Fails by	Giving a wrong answer	Giving a wrong answer <i>and</i> stopping

An agent works in a loop. It thinks → picks an action → sees the result → thinks again. It doesn't stop after one reply; it keeps taking steps until the task is complete or it needs your input.



Why it matters: In Phase 1 you learned the ReAct pattern (Reason + Act) as an advanced *prompting* technique. Claude Code *is* ReAct, built into a product. Everything you learned about clear instructions still applies — it just now drives a tool that can act on those instructions.

A4. Where Claude Code runs

Claude Code is not only a terminal command. The same agent shows up in several places:

- **Terminal / CLI** — the original and most powerful form (`claude` command). This guide focuses here.
- **VS Code & JetBrains extensions** — Claude Code inside your editor, with clickable file links and inline diffs.
- **Desktop app** (Mac & Windows) — a windowed version.
- **Web app** — claude.ai/code, runs in the cloud on your connected repos.

All of them share the same brain and the same core ideas (context, tools, permissions, [CLAUDE.md](#)). Learn the CLI and you know them all.

A5. Key terms glossary

Keep this handy — every later part uses these words.

Term	Plain meaning
Agent	An AI that takes actions in a loop, not just chats
Tool	A capability the agent can use (read a file, run a command, search)
Context window	The agent's short-term memory — everything it can "see" right now
Session	One continuous conversation with Claude Code
Token	A word-piece; both text and code are measured in tokens
CLAUDE.md	A project memo Claude reads automatically every session
Slash command	A built-in command you type starting with <code>/</code> (e.g. <code>/model</code>)
Permission	Your approval before Claude does something risky (edit, run)
MCP	A standard plug for connecting external tools & data sources
Subagent	A helper agent Claude spawns to do part of a job
Headless mode	Running Claude Code with no chat UI, for scripts/automation
Compaction	Summarising old conversation to free up the context window

Part B — Getting Connected

The goal of this part: from nothing installed to Claude editing a file, in five minutes.

B1. Install

Claude Code needs **Node.js** (version 18 or newer). Check with `node --version`. If missing, install it from nodejs.org.

The one-line install:

```
npm install -g @anthropic-ai/claude-code
```

- `npm` is Node's package manager (came with Node).
- `-g` means *global* — install it everywhere, so you can type `claude` in any folder.
- `@anthropic-ai/claude-code` is the official package name.

Verify it worked:

```
claude --version
```


Note: There is also a native installer (a single downloadable binary that doesn't need Node) and the IDE extensions install from their marketplaces. The `npm` route is the most common starting point.

B2. Authentication (connecting your account)

Claude Code has to know *who you are* so it can bill the usage and unlock your models. First run:

```
claude
```

It opens a browser to log in. You have two ways to connect:



- **Subscription (Claude Pro or Max):** log in with your normal Claude account. Usage counts against your plan's limits. Best for individuals — predictable monthly cost.
- **API key (Console):** create a key at console.anthropic.com, paste it in. You pay per token used. Best for automation, teams, and heavy use.

You can switch later with the `/login` and `/logout` slash commands. Your choice is remembered between sessions.

Analogy: A subscription is a gym membership (flat monthly fee, use it as much as the plan allows). An API key is a pay-per-visit pass (you're billed for exactly what you use). Same gym, different billing.

B3. Your first session

Navigate to a project folder and start:

```
cd my-project
claude
```

You're now in an interactive session. Try these, one at a time:

```
> what does this project do?
> explain the folder structure
> find where the login logic lives
> add a comment to the top of the main file explaining what it does
```

Notice the pattern: the first three are **read-only** (Claude just answers). The fourth **changes a file** — Claude will show you the exact edit and ask permission before writing it. You are always in control of changes.

To leave: type `/exit` or press `Ctrl+C` twice.

B4. The interface and modes

Claude Code has three *permission modes* that trade safety for speed. Cycle them with **Shift+Tab**.

Mode	What it does	When to use
Normal (default)	Asks before every edit or command	Learning, unfamiliar code, anything risky
Auto-accept edits	Applies file edits without asking (still asks for commands)	You trust the task and want speed
Plan mode	Read-only — Claude researches and writes a plan, changes <i>nothing</i>	Big tasks: see the plan before any code moves

Plan mode is a superpower for beginners. Ask for something large, let Claude explore and produce a step-by-step plan, review it, then approve. It's the difference between "go rebuild the auth system" (scary) and "show me exactly *how* you'd rebuild it, then we'll go" (safe).

Other interface essentials:


- `Esc` — interrupt Claude mid-action (stop it, redirect it).
- `Ctrl+C` — cancel current input / quit.
- `↑` — recall your previous messages.
- **Drag a file** into the terminal to add it to context.
- **Paste an image** (a screenshot of an error, a design mockup) — Claude can see it.

Part C — How It Works Internally

Understanding the machinery makes you dramatically better at using it. This is the "how it actually works" you asked for.

C1. The agentic loop

Under the hood, every Claude Code task is the same loop repeating:



Walk through a real example — "the login button is broken":

1. **Think:** "I need to find the login button code." → **Act:** search the project for "login".
2. **Observe:** finds `LoginButton.jsx`. → **Think:** "Let me read it." → **Act:** read the file.
3. **Observe:** sees the click handler calls a wrong function name. → **Think:** "That's the bug." → **Act:** edit the file (asks permission).
4. **Observe:** edit applied. → **Think:** "Verify it." → **Act:** run the app / tests.
5. **Observe:** works. → **Done:** reports the fix.

Each **Observe** feeds back into the model's context, so its next decision is informed by what just happened. That feedback is what makes an agent more than a one-shot chatbot.

C2. The context window

The context window is Claude's working memory — every token it can currently see: your messages, the files it has read, command outputs, and `CLAUDE.md`. Modern Claude models have a very large window (hundreds of thousands of tokens — enough for many files at once), and some support up to **1 million tokens**.

Crucial mental model: context is a *desk*, not a *filing cabinet*.

THE DESK (context window)

- Your instructions
- `CLAUDE.md` (project rules)
- Files Claude has opened
- Command outputs it has seen
- The conversation so far

Everything here is "in mind" right now.
What's NOT here, Claude cannot see.

Two consequences you must internalise:

1. **Claude only knows what's on the desk.** It doesn't magically know your whole repo — it *reads files into context* as needed. If it hasn't read a file, it hasn't seen it.
2. **The desk fills up.** Long sessions accumulate a lot of tokens. When it gets full, quality drops and cost rises.

Two fixes:

- `/clear` — wipe the desk and start fresh (do this between unrelated tasks).
- `/compact` — Claude summarises the conversation so far into a short note, freeing space while keeping the gist. Claude Code also *auto-compacts* when the window gets close to full.

Why it matters: The #1 beginner mistake is one giant never-ending session where the desk is buried under ten unrelated tasks. Start fresh (`/clear`) for each new task and Claude stays sharp and cheap.

C3. Tools, the hands of the agent


A model alone can only produce text. **Tools** are what let Claude *act*. Every Claude Code capability is a tool the model chooses to call:

Tool	What it does	Real use
Read	Open a file	"Show me the config"
Edit	Change part of a file	Fix a bug on line 42
Write	Create/overwrite a file	Scaffold a new component
Bash	Run a terminal command	<code>npm test</code> , <code>git status</code>
Grep	Search file <i>contents</i>	Find every use of <code>getUser()</code>
Glob	Find files by name pattern	All <code>*.test.js</code> files
WebFetch / WebSearch	Read the internet	Check library docs
Task	Spawn a subagent	Delegate a big search

How does Claude know *which* tool to use? Each tool has a description — a mini-prompt telling the model what it's for and when to use it. This connects straight to Phase 1's tool-use lesson: **the model only *requests* a tool with structured arguments; the Claude Code program actually *runs* it** and hands back the result. The model never touches your disk directly — that boundary is your safety.

C4. Permissions, the safety layer

Because Claude can edit files and run commands, there's a gate: **permissions**. By default, anything that *changes something* or *runs a command* asks first.



- **Allow once** — yes, this time.

- **Allow always / "don't ask again"** — trust this kind of action for the session (or permanently via settings).
- **Deny** — no. A denied action is *feedback*: Claude should try a different approach, not the same thing again.

You tune this in `.claude/settings.json` with **allow / ask / deny** rules — e.g. auto-allow `npm test` and `git status` (safe, frequent) but always ask for `git push` or anything that deletes. This is how teams make Claude fast without making it dangerous.

Golden rule: the more reversible and read-only an action, the safer it is to auto-allow. Reading and searching — allow freely. Pushing, deleting, deploying — keep the gate.

Part D — Models: Choosing the Right Brain

You asked specifically about models and how using different ones differs. Here's the whole picture.

D1. The model family


Claude Code can run on different Claude models. Think of them as different-sized brains — all smart, trading raw capability against speed and cost.

Model	Personality	Best for
Opus 4.8	The senior architect — deepest reasoning	Hard bugs, big refactors, architecture, tricky logic
Sonnet 5	The reliable senior dev — fast <i>and</i> strong	Everyday coding, the default workhorse
Haiku 4.5	The quick junior — fast & cheap	Simple edits, quick questions, high-volume tasks
Fable 5	Frontier-class alternative in the Claude 5 family	Heavy reasoning workloads

(Model line-ups evolve — always check the current list with `/model` . The idea is stable: bigger = smarter but slower/costlier.)

D2. How they actually differ

The trade-off is a triangle: **capability** ↔ **speed** ↔ **cost**. You can't max all three; you pick per task.



The difference in real terms:

- A **tiny job** ("rename this variable everywhere") — Haiku nails it instantly for a fraction of the cost. Using Opus here is like hiring a principal engineer to fix a typo: wasteful.
- An **everyday feature** ("add form validation") — Sonnet is the sweet spot: strong reasoning, quick, reasonable cost. This is the default for a reason.
- A **gnarly problem** ("this race condition only happens in production") — Opus earns its cost. Its deeper reasoning finds what smaller models miss. Paying more here *saves* money because it solves it in one pass

instead of ten failed cheap attempts.

Why it matters: Matching model to task is the biggest lever on both *quality* and *bill*. Beginners either use the cheapest for everything (and get stuck) or the priciest for everything (and overpay). The pro move: **Sonnet by default, Opus when stuck or for high-stakes design, Haiku for trivial bulk work.**

| D3. Switching models

Interactively, just type:

```
/model
```

It lists available models; pick one. The switch takes effect immediately for the rest of the session.

At launch, pass a flag:

```
claude --model opus
```

Advanced: Claude Code can even use different models for different jobs automatically — e.g. a powerful model for the main reasoning and a lighter one for quick background tasks like summarising. You don't have to manage this, but it's why the tool feels both smart *and* responsive.

| D4. Cost, speed, and fast mode

How billing works depends on how you connected (Part B2):

- **Subscription:** usage draws down your Pro/Max limits — no per-token thinking needed until you hit a cap.
- **API key:** you pay per token — *input* tokens (what Claude reads: your files, context) and *output* tokens (what it writes). Bigger models cost more per token.

Two features that cut cost dramatically:

- **Prompt caching** — reused context (like [CLAUDE.md](#) and files already read) is cached, so you're not billed full price to "re-read" the same thing every turn. This is Phase 1's caching lesson, working automatically.
- **/compact** and **/clear** — smaller context = fewer input tokens = lower cost and faster responses.

Fast mode (**/fast**): on capable models (Opus 4.8/4.7), this speeds up output *without* dropping to a weaker model — you get Opus-quality answers, delivered faster. Toggle it when you want the best brain at higher pace.

Check your spend anytime with **/cost** .

Part E — CLAUDE.md (Project Memory)

The problem it solves: Claude starts each session knowing nothing about *your* project's conventions. Do you use tabs or spaces? Which test command? What's the architecture? Without help, it guesses — and you correct the same things every session.

CLAUDE.md is a memo Claude reads automatically at the start of every session. It's a plain Markdown file in your project root. Put in it the things you'd tell a new teammate on day one.

A good starter CLAUDE.md:

```
# Project: Acme Dashboard

## What this is
A React + TypeScript admin dashboard. API calls live in `src/api/`.

## Commands
- Run tests: `npm test`
- Start dev server: `npm run dev`
- Lint: `npm run lint` (ALWAYS run before committing)


## Conventions
- Use functional components + hooks, never class components.
- 2-space indentation.
- Co-locate a `*.test.tsx` next to every component.

## Don't
- Don't edit anything in `src/generated/` - it's auto-built.
- Don't add new dependencies without asking.
```

Why this is the highest-leverage file you'll write:

- It turns repeated corrections into a permanent rule.
- Every session starts smart instead of clueless.
- It's shared via Git, so your whole team (and their Claude) stays consistent.

Three levels of memory (they stack, most-specific wins):



How to build one:

- Run `/init` in a project — Claude explores the codebase and drafts a `CLAUDE.md` for you. Best starting point.
- Press `#` during a session and type a rule — Claude adds it to memory for you (e.g. `# always use pnpm, not npm`).
- Just edit the file by hand anytime.

Keep it tight. `CLAUDE.md` sits in context *every* session, so every line costs tokens. Write rules, not essays. Prune what's no longer true.

Part F — Commands & Controls

Three kinds of "commands" — don't mix them up.

F1. Slash commands (inside a session)

Typed at the `>` prompt. The ones you'll use constantly:

Command	Does
<code>/help</code>	List everything available
<code>/clear</code>	Wipe context — fresh start for a new task
<code>/compact</code>	Summarise history to free context
<code>/model</code>	Switch model
<code>/init</code>	Generate a <code>CLAUDE.md</code> for this project
<code>/config</code>	Open settings (theme, model, behaviour)
<code>/cost</code>	Show token usage / spend this session
<code>/mcp</code>	Manage MCP connections (Part H)
<code>/agents</code>	Create & manage subagents (Part I)
<code>/review</code>	Review a pull request or the current diff
<code>/login</code> <code>/logout</code>	Switch accounts / auth
<code>/resume</code>	Reopen a past session with its full context
<code>/exit</code>	Quit

Custom slash commands & skills: you can define your own reusable commands as Markdown files in `.claude/commands/`. A `.claude/commands/test.md` becomes `/test`. This packages a repeated instruction ("run the tests, then fix any failures, then lint") into one word.

F2. CLI flags (when launching)

Passed to the `claude` command in your shell:

```
claude --model sonnet      # start on a specific model
claude -p "summarise README" # headless: run one task, print result, exit
claude --resume            # jump back into your last session
claude --continue          # continue the most recent conversation
```

-p (headless / print mode) is the automation gateway. No chat UI — Claude runs the task and prints the answer. This is how you put Claude Code into scripts, Git hooks, and CI pipelines:

```
# Example: auto-generate a commit message from staged changes
git diff --staged | claude -p "write a concise commit message for this diff"
```

F3. Keyboard controls

Key	Action
Shift+Tab	Cycle permission mode (Normal → Auto-accept → Plan)
Esc	Interrupt Claude right now
Ctrl+C	Cancel input / quit
↑ / ↓	Scroll through your message history
#	Quick-add a rule to CLAUDE.md
!	Run a raw bash command yourself (bash mode)

Part G — Using Claude Code Effectively

The "how to use it effectively" you asked for — the habits that separate frustration from flow.

G1. The seven habits

1. **Give it a `CLAUDE.md`.** (Part E.) One file removes 80% of repeated corrections. Do this first, always.
2. **Plan before big changes.** Enter Plan mode (Shift+Tab), let it produce a plan, review, *then* approve. Cheap insurance against a wrong 20-file change.
3. **One task per session.** `/clear` between unrelated jobs. A clean desk = sharper, cheaper Claude (Part C2).
4. **Be specific — Phase 1 still rules.** "Fix the bug" is weak. "The login button does nothing on click; the handler in `LoginButton.jsx` seems to call the wrong function — find and fix it" is strong. Give code + intent + the actual error.
5. **Let it verify its own work.** Ask it to run the tests / run the app after a change. An agent that checks itself catches its own mistakes.
6. **Course-correct early.** If it's going the wrong way, hit **Esc** and redirect. Don't wait for it to finish a wrong path — the context fills with dead ends.
7. **Point it at reality, not your imagination.** Paste the real error, the real screenshot, the real failing test. Claude reasons far better from ground truth than from your paraphrase.

G2. Anti-patterns to avoid

Don't	Instead
One endless session for everything	<code>/clear</code> per task
"Make it better" (vague)	Name the file, the symptom, the goal
Opus for every trivial edit	Sonnet default, Opus when stuck
Blindly approving every edit	Read the diff — you're the reviewer
Ignoring <code>CLAUDE.md</code>	Invest 10 min once, save hours
Letting context bloat silently	Watch <code>/cost</code> , <code>/compact</code> when big

G3. The reviewer mindset

Claude Code makes *you* the senior reviewer, not the typist. Your job shifts from *writing* every line to *deciding* whether the written lines are right. That's a promotion — but only if you actually read the diffs. Approval is a responsibility, not a formality.

Part H — MCP (Model Context Protocol)


H1. The problem MCP solves

Claude Code can read your files and run commands. But your work also lives *elsewhere*: a database, a Jira board, Figma designs, Slack, your company's internal API. How does Claude reach those?

Without a standard, every tool would need a custom, hand-built integration — an N×M mess.

MCP is a universal adapter. It's an open standard (created by Anthropic) that defines *one* way for AI tools to connect to external systems. Build a connector once, and any MCP-aware app can use it.

Analogy: MCP is the **USB-C of AI tools**. Before USB-C, every device had its own charger. USB-C standardised the plug — one cable, any device. MCP standardises how AI connects to data and tools — one protocol, any source.



H2. How it works

An **MCP server** is a small program that exposes a system (say, your database) as a set of **tools** Claude can call — e.g. `query_database`, `list_tables`. Claude Code is the **MCP client** that connects to those servers. Once connected, those external tools appear alongside Claude's built-in ones, and it uses them in the same agentic loop (Part C1).

Servers can offer three things: **tools** (actions Claude can take), **resources** (data Claude can read), and **prompts** (ready-made prompt templates).

H3. Connecting a server

```
# Add an MCP server (example shape – each server has its own command)
claude mcp add postgres-db -- npx -y @some/postgres-mcp-server

# See what's connected
claude mcp list
```

Inside a session, manage connections with `/mcp`. Connections can be scoped to just you, shared with your project (via a checked-in config), or global.

| H4. Real-world example

Task: "Which users signed up last week but never logged in?"

- **Without MCP:** you open a SQL client, write the query, run it, copy results back to Claude. Manual.
- **With a Postgres MCP server:** Claude calls `query_database` itself, gets the rows, and answers directly — then can even draft the follow-up email. The database became one of Claude's hands.

Why it matters: MCP is how Claude Code graduates from "coding assistant" to "assistant for your whole workflow" — reaching the real systems where your work actually lives. **Safety note:** only connect servers you trust, and be cautious granting write access to production systems.


Part I — Agents & Workflows

The most advanced part — how one Claude becomes a coordinated team.

11. Subagents: delegation

A big task can blow up the main context: reading 30 files to answer one question buries everything else on the desk. The fix: **subagents**.

A subagent is a *separate* Claude with its *own* fresh context, spawned to do one focused job. It does the work, and returns only the *conclusion* to the main agent — not the 30 files it read.



Two big wins:

- **Context stays clean.** The main desk never sees the 30 files — only the tidy result.
- **Parallelism.** Independent subagents run *at the same time*. Three searches that would take 3× sequentially finish in roughly 1×.

Custom subagents: define specialists as Markdown files in `.claude/agents/` (or via `/agents`) — e.g. a `code-reviewer` agent with a strict review prompt, or a `test-writer` agent. Each has its own instructions, its own allowed tools, and can even be pinned to a specific model (a cheap Haiku reviewer, an Opus architect).

12. Hooks: automation

Hooks are commands the *harness* (not Claude) runs automatically at set moments — configured in `.claude/settings.json`. They're deterministic: they always fire, regardless of what Claude decides.

Common trigger points: `PreToolUse` (before a tool runs), `PostToolUse` (after), `Stop` (when Claude finishes).


Real examples:

- **After every file edit**, auto-run the code formatter → your code is always clean, no reminders.
- **Before any git push**, run the test suite → never push broken code.
- **When Claude finishes**, play a sound → you can look away and get pinged when it's done.

Why hooks vs. just asking Claude? Because "always" needs a guarantee. Asking Claude "please format after editing" works *usually*; a hook works *every single time*. Rules you can't afford to have skipped belong in hooks.

13. Orchestration: the full workflow

Put it together and Claude Code becomes an orchestrator:



For very large jobs (a migration across 100 files, a codebase-wide audit), you can fan work out to *many* agents in parallel — each handling one slice, with a verification pass checking the results. This is how Claude Code tackles work too big for a single context to hold. Start simple (one agent, plan mode); reach for orchestration only when the task genuinely demands scale.

Part J — Real-World Scenarios

Concrete "when would I actually use this" walkthroughs. Each maps to what you learned.

Scenario 1 — Onboarding to an unfamiliar codebase

You joined a project with 500 files and no docs.

```
> Give me a tour: what does this app do, what's the tech stack,  
and where does a request flow from URL to database?
```

Claude reads the key files and explains. Follow with `/init` to capture what it learned into [CLAUDE.md](#). **Uses:** read tools, context, [CLAUDE.md](#).

Scenario 2 — Fixing a bug from an error message

```
> Running `npm test` gives this error: [paste full error].  
Find the cause and fix it, then re-run the tests to confirm.
```

Claude searches, reads, edits, and *verifies by running the tests itself* — the full agentic loop. **Uses:** the loop, Bash tool, self-verification (G1 #5).

Scenario 3 — Building a feature safely

```
> [Shift+Tab into Plan mode]  
> Add CSV export to the reports page. Show me your plan first.
```

Review the plan, approve, let it build, read the diff. **Uses:** Plan mode, reviewer mindset.

Scenario 4 — A large, risky refactor

```
> Rename `getUserData` to `fetchUserProfile` across the whole codebase,  
update all call sites, and make sure nothing breaks.
```

Best on **Opus**, in Plan mode first. Claude finds every usage, edits consistently, runs tests. **Uses:** model choice (D2), Grep, verification.

Scenario 5 — Data question via MCP

```
> How many orders were refunded last month, by product category?
```

With a database MCP server connected, Claude queries directly and answers — no manual SQL. **Uses:** MCP (Part H).

Scenario 6 — Automating a repetitive chore (headless)

```
# In a Git pre-commit hook or a nightly script:  
claude -p "review the staged diff for security issues; list any found"
```

No UI — Claude runs, prints, exits. **Uses:** headless mode (F2), automation.

Scenario 7 — Learning while building (ties to your goals)

```
> Explain this authentication code to me like Feynman would –  
simple language, an analogy, then quiz me on it.
```

Claude Code isn't only for shipping; it's a tutor sitting inside real code — exactly the active-recall style from Phase 1.

Part K — Exercises & Deliverables

Do these in order. Each builds a real skill. Save your work — these are your Phase 2 proof.

1. **Install & connect.** Get `claude --version` working and log in. Deliverable: a screenshot of your first successful session.
 2. **First edits.** In a small practice repo, ask Claude to add a feature and a comment. Practice the permission prompts (allow / deny). Deliverable: the diff it produced.
 3. **Write a `CLAUDE.md`.** Run `/init`, then hand-tune it with 5 real rules for the project. Deliverable: the file.
 4. **Model comparison.** Give the *same* medium task to Haiku, Sonnet, and Opus. Note speed, quality, and (if on API) cost. Deliverable: a 5-line comparison table — this is where "how models differ" becomes real to you.
 5. **Plan mode.** Take a bigger task, run it through Plan mode, and screenshot the plan *before* approving.
 6. **Custom command.** Create a `.claude/commands/` file that becomes a `/`-command you use. Deliverable: the command file.
 7. **Connect one MCP server.** A simple one (filesystem, or a public API server). Deliverable: `claude mcp list` output showing it connected.
 8. **Subagent task.** Ask Claude to spawn agents to research two independent questions about a repo in parallel. Notice the speed.
 9. **Headless automation.** Write a one-line `claude -p` command that does something useful in your shell (e.g. summarise a file). Deliverable: the command + its output.
 10. **Reflection.** In your own words (Feynman test): explain the agentic loop, the context window, and when you'd pick each model — no notes.
-

Part L — Q&A for Every Topic

Foundations

Q: In one sentence, how is Claude Code different from the Claude chat app? A: The chat app talks; Claude Code *acts* — it reads, edits, and runs things inside your actual project, in a loop, until the job is done.

Q: What makes it an "agent" and not just a chatbot? A: The loop. It thinks → acts with a tool → observes the result → thinks again, repeating until the goal is met, instead of stopping after one reply.

Q: Do I need to be a terminal expert? A: No. You type plain English. The terminal is just the doorway; the conversation is normal language.

Connecting & Interface

Q: What are the two ways to authenticate, and who's each for? A: A Claude Pro/Max **subscription** (flat cost, best for individuals) or an **API key** from the Console (pay-per-token, best for automation/teams).

Q: What is Plan mode and why is it a beginner's best friend? A: A read-only mode where Claude researches and writes a step-by-step plan without changing anything — so you approve the approach *before* any code moves.

Internals

Q: What is the context window, in one image? A: A desk. Everything Claude can see right now sits on it — your messages, opened files, [CLAUDE.md](#), command outputs. What's not on the desk, Claude can't see. When it fills, use `/clear` or `/compact`.

Q: Who actually runs the tools — the model or the program? A: The program. The model only *requests* a tool with arguments; Claude Code executes it and returns the result. That boundary is your safety.

Q: Why does `/clear` between tasks matter? A: A clean context is sharper and cheaper. One endless session buries the current task under old, irrelevant tokens.

Models

Q: Give the one-line rule for picking a model. A: Sonnet by default, Opus when stuck or for high-stakes design, Haiku for trivial bulk edits.

Q: Why can paying for Opus be *cheaper* overall? A: It can solve a hard problem in one pass instead of ten failed cheap attempts — fewer tokens and less of your time wasted.

Q: What does fast mode do — is it a weaker model? A: No. On Opus 4.8/4.7 it speeds up output while keeping Opus-level quality; it does not downgrade the brain.

CLAUDE.md

Q: What is [CLAUDE.md](#) and why is it the highest-leverage file? A: A Markdown memo Claude auto-reads every session. It turns repeated corrections into permanent rules, so every session starts smart and your team

stays consistent.

Q: Fastest way to create one? A: Run `/init` — Claude explores the repo and drafts it. Then trim it to real rules; keep it short because it costs context every session.

| Commands

Q: Difference between a slash command and a CLI flag? A: A slash command (`/model`) runs *inside* a session; a CLI flag (`--model`) is passed to `claude` when you launch it in the shell.

Q: What is headless mode for? A: `claude -p "task"` runs one task with no chat UI and prints the result — the gateway to scripts, Git hooks, and CI automation.

| MCP

Q: What is MCP in one analogy? A: The USB-C of AI tools — one open standard to plug Claude into any external system (databases, Jira, Figma, your API) instead of custom one-off integrations.

Q: After connecting an MCP server, what changes for Claude? A: The server's actions appear as new tools Claude can call in its normal loop — e.g. it can query your database itself instead of you doing it by hand.

| Agents & Workflows

Q: Why use a subagent instead of doing everything in the main session? A: It keeps the main context clean (returns only the answer, not the 30 files it read) and lets independent work run in parallel.

Q: Hooks vs. just instructing Claude — when do you need a hook? A: When you need "always." Asking Claude works *usually*; a hook (in settings.json) fires *every time*, deterministically — e.g. run tests before every push.


| Effective Use

Q: The single mental model to carry from Phase 1 into Phase 2? A: Still "briefing a brilliant new employee on day one" — but now that employee has hands. Clear briefing + a good `CLAUDE.md` + reviewing the work = the whole craft.

Q: What's the biggest beginner mistake? A: One giant never-ending session for ten unrelated tasks. Clear between tasks, be specific, and actually read the diffs.

Part M — Interview Preparation

Now flip roles. This part drills the same material as **interview questions** — the way a senior HR + technical panel actually asks: what, why, differences, scenarios, and behaviour. Every answer uses one shape you can say aloud: **definition** → **why** → **example** → **(trade-off if asked)**. The 🌀 notes reveal the real thing being tested behind each question.



Round 1 — Fundamentals (What & Why)

Q1. What is Claude Code, in one sentence? A: Claude Code is an agentic command-line tool that puts Claude *inside* your project — it reads your files, runs your commands, and edits your code in a loop, while you supervise. It turns Claude from something you *chat with* into something you *delegate to*. 🌀 *Interviewer is checking:* can you compress a concept into one clean sentence.

Q2. Why does Claude Code exist — wasn't the chat app enough? A: A chat model only sees the text you paste, so *you* do all the mechanical work — finding files, applying edits, running tests. Claude Code removes that gap by giving Claude the same tools a developer uses (filesystem, terminal, search). The "why" is leverage: you stop being the typist and become the reviewer. 🌀 *Interviewer is checking:* do you understand the problem it solves, not just the feature.

Q3. How is an "agent" different from an "assistant"? A: An assistant answers once and stops. An agent runs a loop — think → act with a tool → observe the result → think again — repeating until the goal is met. The loop, plus the ability to take real actions, is the whole difference. Claude Code is essentially the ReAct pattern built into a product. 🌀 *Interviewer is checking:* the loop concept — the heart of everything.

Q4. Give a real example where Claude Code beats a chatbot. A: "The login button does nothing on click." A chatbot explains possible causes and I implement the fix by hand. Claude Code *searches* the repo, *reads* the handler, spots that it calls the wrong function, *edits* the file, and *runs the tests* to confirm — end to end, one instruction. The intelligence is similar; the reach is not. 🌀 *Interviewer is checking:* can you ground theory in a concrete story.

Q5. Do you need to be a terminal expert to use it? A: No. You type plain English — "fix the login bug" is a valid command. The terminal is just the doorway; the conversation inside is normal language. That's a deliberate design choice to lower the barrier.

Round 2 — How It Works Internally

Q6. Walk me through the agentic loop. A: Four repeating steps. **Think** — the model decides the next action. **Act** — it calls a tool (read, edit, bash, search), asking permission if the action changes something. **Observe** — the result is fed back into its context. **Repeat** — informed by that result, it decides again, until the goal is met or it needs me. Each observation makes the next decision smarter — that feedback is what makes it an agent, not a one-shot reply.

Q7. What is the context window, and why should I care? A: It's Claude's working memory — every token it can currently see: my messages, files it has read, command outputs, and `CLAUDE.md`. The mental model is a **desk, not a filing cabinet**: Claude only knows what's on the desk right now; a file it hasn't read, it hasn't seen. I care because (a) it explains why I must point it at the right files, and (b) the desk fills up on long sessions, which lowers quality and raises cost — so I `/clear` between tasks and `/compact` when it gets big. 🧠 *Interviewer is checking:* the "desk" model and the practical habits that follow from it.

Q8. Who actually runs the tools — the model or the program? A: The program. The model only *requests* a tool with structured arguments; the Claude Code harness executes it and returns the result. The model never touches my disk directly. That boundary is exactly why permissions can exist — there's a gate between "the AI wants to" and "it happened."

Q9. How does Claude decide which tool to use? A: Each tool has a description that acts like a mini-prompt telling the model what it's for and when to use it. Given the goal and what it has observed, the model picks the tool whose description best fits the next step — the same tool-use principle from prompt engineering, applied automatically.

Q10. How do permissions work, and why do they matter? A: Anything that changes something or runs a command asks first: **allow once**, **allow always**, or **deny**. A denial is *feedback* — Claude should try another approach, not repeat the same call. I tune it in `settings.json` with allow/ask/deny rules: auto-allow safe, frequent, read-only actions (`git status` , `npm test`), always gate irreversible ones (`git push` , deletes, deploys). The rule of thumb: the more reversible and read-only, the safer to auto-allow.

| Round 3 — Models & Cost Judgement

Q11. Claude Code can run on different models — how do you choose? A: It's a trade-off triangle: **capability** ↔ **speed** ↔ **cost**. My default is **Sonnet** — strong and fast for everyday work. I reach for **Opus** when I'm stuck on a hard bug, doing a big refactor, or making a high-stakes design decision — its deeper reasoning earns its cost. I drop to **Haiku** for trivial, high-volume edits like a bulk rename, where paying for Opus is like hiring a principal engineer to fix a typo. 🧠 *Interviewer is checking:* judgement and cost-awareness, not memorised model names.

Q12. Why can using the *more expensive* model be *cheaper* overall? A: Because a stronger model can solve a hard problem in one pass instead of ten failed cheap attempts. Ten Haiku retries that never land can burn more tokens — and far more of my time — than one Opus pass that solves it. Cost isn't per-token; it's per-solved-problem.

Q13. What's "fast mode" — is it a weaker model? A: No, and that's a common trap in the question. On Opus 4.8/4.7, fast mode speeds up output while keeping Opus-level quality — same brain, quicker delivery. It does not downgrade to a smaller model.

Q14. Subscription vs API key — what's the difference and who's each for? A: A **Claude Pro/Max subscription** is a flat monthly cost with usage limits — best for individuals; predictable. An **API key** from the Console is pay-per-token (input + output) — best for automation, teams, and heavy use. Same tool, different billing model — like a gym membership versus a pay-per-visit pass.

| Round 4 — CLAUDE.md, Commands & Config

Q15. What is `CLAUDE.md` and why do you call it the highest-leverage file? A: It's a Markdown memo Claude reads automatically at the start of every session — the things you'd tell a new teammate on day one: what the

project is, the test command, conventions, what not to touch. It's highest-leverage because it turns repeated corrections into permanent rules: every session starts smart instead of clueless, and because it's committed to Git, the whole team's Claude stays consistent.

Q16. How do you create a good one, and what's the #1 mistake? A: Fastest start is `/init` — Claude explores the repo and drafts it; then I trim to real rules. I can also press `#` mid-session to append a rule. The #1 mistake is bloat: it sits in context *every* session, so every line costs tokens. Write rules, not essays, and prune what's no longer true.

Q17. Difference between a slash command and a CLI flag? A: A slash command (`/model` , `/clear`) runs *inside* a live session. A CLI flag (`--model` , `-p`) is passed to the `claude` command when I launch it from the shell. Different moments: one steers an ongoing conversation, the other configures the start.

Q18. `/clear` vs `/compact` — when do you use each? A: `/clear` wipes the context for a genuinely new, unrelated task — a clean desk. `/compact` summarises the current conversation to free space while *keeping the gist* — I use it when I'm deep in one long task that's filling the window but I still need its history. Rule: new task → clear; same task, too big → compact.

Q19. What's headless mode and why does it matter? A: `claude -p "task"` runs one task with no chat UI and prints the result. It matters because it's the automation gateway — I can drop Claude Code into shell scripts, Git hooks, and CI pipelines. Example: pipe a staged diff into `claude -p "write a commit message"` to auto-generate commit messages.

Round 5 — MCP & Integrations

Q20. What is MCP, in one analogy? A: MCP (Model Context Protocol) is the **USB-C of AI tools** — one open standard to plug Claude into any external system (databases, Jira, Figma, Slack, your internal API). Before it, every integration was a custom one-off; MCP means you build a connector once and any MCP-aware app can use it.

Q21. How does it actually work? A: An MCP *server* exposes a system as a set of tools (e.g. `query_database` , `list_tables`). Claude Code is the MCP *client* that connects to it. Once connected, those external tools appear alongside Claude's built-in ones, and it uses them in the same agentic loop. Servers can offer tools (actions), resources (data to read), and prompts (templates).

Q22. Give a real example of MCP earning its keep. A: "Which users signed up last week but never logged in?" Without MCP I open a SQL client, write the query, run it, paste results back. With a Postgres MCP server, Claude calls `query_database` itself, gets the rows, answers, and can even draft the follow-up email. The database became one of Claude's hands. 🚀 *Interviewer is checking:* can you show the before/after value, plus awareness that you only connect servers you trust — especially with write access to production.

Round 6 — Agents, Hooks & Workflows

Q23. Why use a subagent instead of doing everything in the main session? A: Two reasons. **Clean context** — a subagent has its own fresh window, does a focused job (say, read 30 files to find every API endpoint), and returns only the *conclusion* to the main agent, so the main desk never gets buried. **Parallelism** — independent subagents run at the same time, so three searches finish in roughly the time of one.

Q24. Hooks vs. just instructing Claude — when do you need a hook? A: When you need **"always."** Asking Claude "please format after editing" works *usually*; a hook in `settings.json` fires *every single time*,

deterministically, because the harness runs it, not the model. Anything you can't afford to have skipped — run tests before every push, format after every edit — belongs in a hook.

Q25. Describe orchestrating a non-trivial feature end to end. A: Start in **Plan mode** so Claude breaks the task down read-only and I approve the approach. Then it implements — possibly fanning independent parts to subagents in parallel. A **hook** runs the tests automatically after edits. Finally I read the diff and approve. Simple tasks stay single-agent; I only reach for heavy orchestration when the work genuinely needs scale, like a 100-file migration.

Round 7 — Product Differences (the comparison round)

This round tests whether you understand the whole Claude ecosystem, not just one tool. These are the "differences" questions interviewers love.

Q26. Difference between Claude (chat), Claude Code, and Claude Cowork? A: Same underlying model and agentic foundation — three different jobs and audiences.

	Claude (chat)	Claude Code	Claude Cowork
Who	Anyone	Developers	Knowledge workers (non-coders)
Where	Browser / app	Terminal & IDE	Claude desktop app (+ web/mobile)
Does	Advises — produces text/artifacts	Reads, edits, runs code locally	Works in your files & connected apps (Slack, Drive)
Style	Conversation	Power, control, extensibility	Simple, safe, no command line
Setup	None	Install & configure yourself	None — just start chatting
Runs in	The chat	Your real machine	An isolated/sandboxed VM

One-liner: Claude chat *advises*; Cowork is a *general coworker* that does office/file tasks for non-coders in a safe sandbox; Claude Code is the *power tool* for engineers with full control and extensibility. Anthropic's own guidance: start with Cowork to feel what's possible, graduate to Claude Code when you want precision, speed, and control. 🌀 *Interviewer is checking:* that you know they're the *same engine, different jobs* — and can match each to a user.

Q27. So when would you pick Cowork over Claude Code? A: Cowork when the work is **non-coding knowledge work** — turning research into a deck, building a spreadsheet, prepping meeting briefs from Slack + Drive — and the user doesn't want a terminal. It's simpler and safer (sandboxed, isolated from the wider internet) and can run tasks in the background even when your device is offline. Claude Code when the work is **software** and I want precision, speed, extensibility (MCP, hooks, subagents), and reliability on long multi-step tasks. Cowork can stall on very long workflows where Code holds up.

Q28. Difference between Claude Artifacts and Claude Design? A: **Artifacts** are the general *rendering layer* — when Claude produces substantial standalone output (an HTML app, code, an SVG, a doc, a diagram), it opens in a dedicated side window instead of being buried in chat, and it's interactive and shareable. It's for *any* kind of output.

Claude Design is a *specialised product* (Anthropic Labs, powered by Opus 4.7) focused purely on **visual design work** — prototypes, wireframes, slides, one-pagers, pitch decks. Its differentiators: it can read your codebase and Figma files to extract your **design system**, then apply your brand (colours, typography, components) automatically to every project; you refine via inline comments, direct edits, and sliders; and you export to Canva, PDF, PPTX, or HTML — or hand the whole thing to Claude Code as a build bundle.

One-liner: Artifacts = a general output canvas *inside chat*; Claude Design = a dedicated *design studio* that's design-system-aware and export/handoff-ready. 🌀 *Interviewer is checking:* general-purpose vs specialised, and the design-system + handoff angle.

Q29. Claude Code (CLI) vs the IDE extension — same thing? A: Same brain and same concepts (context, tools, permissions, [CLAUDE.md](#), MCP). The difference is surface: the CLI lives in the terminal (most scriptable, headless-capable); the IDE extension embeds it in VS Code / JetBrains with clickable file links and inline diffs. Choose by workflow, not capability.

Q30. Claude Code vs a tab-completion tool like GitHub Copilot? A: Different category. Copilot-style tools *autocomplete* as you type — you're still driving line by line. Claude Code is an *agent* — you give a goal and it plans, edits across many files, runs commands, and verifies. Autocomplete helps you write the next line; an agent completes the whole task. They're complementary.

Q31. Plan mode vs Auto-accept mode — when each? A: **Plan mode** is read-only: Claude researches and writes a step-by-step plan, changing nothing, so I approve the approach before any code moves — ideal for big or risky tasks. **Auto-accept** applies edits without asking each time — ideal when I already trust the task and want speed. They sit at opposite ends of the safety-vs-speed dial; Shift+Tab cycles between them.

Round 8 — Scenario / Situational (STAR)

Answer these with Situation → Task → Action → Result. The interviewer wants your judgement under realistic conditions.

Q32. Scenario: You join a 500-file codebase with no documentation. How do you use Claude Code on day one? A: I open it in the repo and ask for a tour — what the app does, the stack, how a request flows from URL to database. Claude reads the key files and explains. Then I run `/init` to capture what it learned into a [CLAUDE.md](#) so every future session starts informed. *Result:* I'm productive in an hour instead of a week, and the team benefits from the memo I committed.

Q33. Scenario: A test is failing in CI and you can't reproduce why. Walk me through it. A: I paste the *full* error into Claude Code and ask it to find the cause, fix it, and re-run the tests to confirm. It searches, reads the relevant code, forms a hypothesis, edits, and verifies by running the suite itself — the full agentic loop with self-verification. I read the diff before approving. *Key move:* I give it ground truth (the real error), not my paraphrase.

Q34. Scenario: You must rename a core function used across 80 files without breaking anything. What's your approach — and which model? A: This is high-stakes, so **Opus**, and I start in **Plan mode**. Claude finds every call site (Grep), shows me the plan, then edits consistently and runs the tests. I review the diff. *Why these choices:* the blast radius is large, so I want the deepest reasoning and a plan I approve before anything changes.

Q35. Scenario: A junior teammate runs one endless Claude Code session for the whole day and complains it's "getting dumber and expensive." What's happening and what do you advise? A: Classic context bloat — the desk is buried under ten unrelated tasks, so quality drops and input-token cost climbs. Advice: `/clear`

between unrelated tasks, `/compact` when a single long task fills the window, and write a `CLAUDE.md` so it stops re-learning basics. *Result:* sharper answers, lower bill. 🧠 *Interviewer is checking:* can you diagnose from symptoms and coach a teammate.

Q36. Scenario: The team wants Claude to *never* push code without running tests first. Memory? A polite instruction? How do you guarantee it? A: Neither instruction nor memory *guarantees* it — those work "usually." I'd add a **hook** in `settings.json` on the pre-push / pre-tool event that runs the test suite, because hooks are deterministic and fire every time regardless of what the model decides. Guarantees belong in hooks, not prompts.

Q37. Scenario: A non-technical marketing colleague wants AI help building a campaign deck from scattered docs and Slack threads. Do you set them up with Claude Code? A: No — that's a **Cowork** job, not Claude Code. They don't want a terminal, the work is non-coding knowledge work across files and connected apps (Slack, Drive), and Cowork's sandbox is simpler and safer for them. I'd reserve Claude Code for the engineers. *Judgement:* match the tool to the person and the task.

Q38. Scenario: You need Claude to answer questions from the production database directly. How, and what's your safety stance? A: Connect a database **MCP server** so Claude can query it in its loop instead of me hand-running SQL. Safety-wise, I only connect servers I trust, scope it read-only for production, and keep write/destructive access gated behind explicit permission. Convenience never overrides guarding production.

Round 9 — Behavioural & Judgement (the HR round)

Softer questions, but they reveal how you actually work with an AI agent. Be honest and specific.

Q39. How do you make sure AI-written code is trustworthy? A: I treat Claude Code as promoting me to *reviewer*. I read every diff, make it verify its own work by running tests or the app, keep risky actions gated behind permissions, and lean on Plan mode for anything large. Approval is a responsibility, not a formality — the moment I rubber-stamp diffs, the value collapses.

Q40. Tell me about a time a prompt/task went wrong. What did you do? A: (STAR.) When Claude heads down the wrong path, I don't wait for it to finish — I hit **Esc**, redirect with more specifics (the exact file, the real symptom, the constraint I forgot to state), and often switch to Plan mode so I can see its intent before it acts. Course-correcting early keeps the context clean of dead ends. 🧠 *Interviewer is checking:* that you steer actively rather than passively accept output.

Q41. What makes someone *good* at using Claude Code versus frustrated by it? A: The good ones invest in a `CLAUDE.md`, plan big changes, keep one task per session, are specific (code + intent + real error), let it verify itself, and actually read the diffs. The frustrated ones run one giant vague session, use the wrong model for the job, and approve blindly. Same tool, opposite outcomes — it's habits, not luck.

Q42. How do you keep costs under control on a team? A: Default to Sonnet and escalate to Opus only when justified; `/clear` and `/compact` to keep context small; lean on prompt caching (reused context like `CLAUDE.md` isn't re-billed at full price); watch `/cost`; and use Haiku for bulk trivial work. Cost discipline is mostly context discipline.

Q43. Where would you *not* use an AI agent? A: Where the action is irreversible and unverifiable, or where I can't review the outcome — I keep those gated or manual. And I never assume it's right on facts it can't check: hallucination never fully goes away, so trust-but-verify is permanent, especially for anything touching production or security.

Round 10 — Rapid-Fire One-Liners

Short questions, one-breath answers — practice saying these fast and clean.


- **What makes Claude Code an agent?** The loop: think → act → observe → repeat.
- **Context window in one word?** Memory — the "desk" of what it can see now.
- **Who runs the tools?** The program, not the model — the model only asks.
- **Default model?** Sonnet; Opus when stuck; Haiku for trivial bulk.
- **Highest-leverage file?** `CLAUDE.md`.
- **Fastest way to draft it?** `/init`.
- `/clear` vs `/compact` ? New task vs same task too big.
- **MCP in three words?** USB-C for AI.
- **Why a subagent?** Clean context + parallelism.
- **Why a hook not a prompt?** Hooks guarantee "always."
- **Cowork vs Code?** Same engine — knowledge work vs software.
- **Artifacts vs Design?** General output canvas vs specialised design studio.
- **Headless mode?** `claude -p` for scripts and CI.
- **Biggest beginner mistake?** One endless session; approving diffs blindly.

The comparison cheat-sheet

If you memorise one table for the "differences" round, make it this.

Compare	Key distinction
Assistant vs Agent	Answers once vs acts in a loop until done
Claude vs Claude Code	Advises in chat vs edits & runs code locally
Claude Code vs Cowork	Developer/terminal/power vs knowledge-worker/desktop/sandboxed — same engine
Artifacts vs Claude Design	General output canvas in chat vs specialised, design-system-aware design studio
CLI vs IDE extension	Same brain; terminal (scriptable) vs editor (inline diffs)
Claude Code vs Copilot	Goal-driven agent vs line-by-line autocomplete
Opus vs Sonnet vs Haiku	Deepest vs balanced default vs fast & cheap
Plan vs Auto-accept	Approve the plan first vs apply edits fast
Subagent vs main agent	Fresh isolated context, runs in parallel vs the main session
Hook vs instruction	Deterministic "always" vs "usually"
Slash command vs CLI flag	Inside a session vs at launch
<code>/clear</code> vs <code>/compact</code>	Wipe for a new task vs summarise to free space
Subscription vs API key	Flat plan (individuals) vs pay-per-token (teams/automation)
Fast mode vs smaller model	Same Opus quality, quicker vs actually weaker brain

Closing: how to answer any interview question



Interview recall card Structure every answer: definition → why → example (→ trade-off if asked). **The loop** is the heart of "agent": think → act → observe → repeat. **The desk** is the context window: it only sees what's on it. **Model judgement:** Sonnet default · Opus when stuck / high-stakes · Haiku for trivial bulk. **CLAUDE.md** = highest-leverage file; **hooks** = deterministic "always"; **MCP** = USB-C for AI. **The big three:** Claude chat *advises* · Cowork *does knowledge work (sandboxed)* · Claude Code *does software (power)*. **Artifacts vs Design:** general output canvas vs specialised, design-system-aware studio. **Always:** you're the reviewer — read the diffs, make it verify, trust but verify.

Part N — Resources & Quick Reference

Learn more (free, official)

- **Claude Code docs:** <https://docs.anthropic.com/en/docs/claude-code/overview>
- **Claude Code quickstart:** <https://docs.anthropic.com/en/docs/claude-code/quickstart>
- **CLAUDE.md & memory:** <https://docs.anthropic.com/en/docs/claude-code/memory>
- **MCP (Model Context Protocol):** <https://modelcontextprotocol.io>
- **Subagents & hooks:** <https://docs.anthropic.com/en/docs/claude-code/sub-agents>
- **Anthropic on YouTube:** <https://www.youtube.com/@AnthropicAI>

The essential commands card

Start: `claude` · **Headless:** `claude -p "task"` · **Model at launch:** `claude --model sonnet` **In-session:**
`/help` `/clear` `/compact` `/model` `/init` `/cost` `/mcp` `/agents` `/review` `/resume` `/exit` **Keys:**
Shift+Tab = permission mode · Esc = interrupt · # = add memory · ! = bash · ↑ = history


One-page recall card

Claude Code = Claude *inside* your project — reads, edits, runs, in a loop. An **agent**, not a chatbot. **Agentic loop** = think → act (tool) → observe → repeat → done. **Context window** = the desk; only what's on it is seen. `/clear` per task, `/compact` when full. **Tools** = Read · Edit · Write · Bash · Grep · Glob · Web · Task. The *program* runs them; the model only *asks*. **Permissions** = allow / ask / deny. Read freely; gate push/delete/deploy. **Models** = Sonnet default · Opus when stuck / high-stakes · Haiku for trivial bulk. Capability ↔ speed ↔ cost. **CLAUDE.md** = auto-read project memo; highest-leverage file; `/init` to draft, `#` to add, keep it short. **Commands** = slash (in-session) vs CLI flags (launch) vs custom `/`-commands in `.claude/commands/`. **MCP** = USB-C for AI; plug in databases/Jira/Figma; servers expose tools Claude calls. **Subagents** = fresh-context helpers, run in parallel, return only the answer. **Hooks** = deterministic automation in settings.json — for things that must happen *every* time. **Effective:** **CLAUDE.md** first · plan big changes · one task per session · be specific · make it verify · read the diffs.

*Chat models answer. Claude Code does. Brief it well, give it a **CLAUDE.md**, pick the right model, review its work — and you stop typing code and start directing it. That is the whole craft of Phase 2.*

Part 0 — The Big Picture: How It All Connects

Everything in this guide, in one map. Read it top to bottom as the life of a single task — from your goal to the finished work.



Reading the map, piece by piece:

- **You** → **CLI** — you type a goal in plain English into Claude Code (terminal, IDE, desktop, or web). Same brain everywhere.
- **The Session** — one continuous conversation. Everything lives inside the **context window** — the agent's "desk" of what it can see right now.
- **Tokens** — the unit that both *fills* the context window and *measures the bill*. Bigger context = more tokens = slower and costlier, which is why you `/clear` and `/compact`.
- **CLAUDE.md** — the project memo loaded automatically at the start, so the agent begins every session already knowing your rules.
- **Main Agent + the loop** — Claude runs *think* → *act* → *observe* → *repeat* until the goal is met. This loop is what makes it an agent, not a chatbot.
- **Model** — the brain powering the agent. Swap between Opus, Sonnet, and Haiku to trade capability against speed and cost.
- **Tools** — the agent's built-in hands: read, edit, write, run commands, search, browse the web.
- **Skills** — packaged, reusable instructions (and custom slash commands) that turn a repeated procedure into one trigger.
- **MCP** → **Connectors** — the universal plug that extends the agent's reach to *external* systems (databases, Jira, Figma, Slack, Drive, your own API), which then appear as extra tools.

- **Sub-agents** — helper agents with their own fresh context that do focused work in parallel and hand back only the answer, keeping the main desk clean.
- **Workflows** — orchestration that fans work out to many agents and verifies the results, for jobs too big for one context.
- **Overall Work** — the finished output: edited code, fixed tests, data answers, deliverables — with *you* as the reviewer approving the diff.

The whole system in one line: *You brief the agent (CLI + [CLAUDE.md](#)); it thinks and acts in a loop (context + tokens), powered by a model, using tools, skills, MCP connectors, sub-agents, and workflows — and produces work you review.*